

Artificial Intelligence Answer Sheet - Tejas (Student 39)

Question 1: Explain the differences between supervised, unsupervised, and reinforcement learning in machine learning. Provide examples of applications for each approach.

Answer:

Supervised learning is a paradigm where algorithms are trained on datasets containing both inputs and their corresponding outputs. The model learns to map inputs to outputs, making it highly effective for prediction tasks. This approach is widely used in classification (such as spam detection or tumor identification in medical scans) and regression (like forecasting house prices or stock values). Supervised learning is most powerful when there is an abundance of labeled historical data, enabling applications in credit scoring, sentiment analysis, and medical diagnostics.

Unsupervised learning, by contrast, works with datasets that lack explicit output labels. The algorithm must autonomously uncover patterns, groupings, or structures within the data. Techniques such as clustering (e.g., customer segmentation for marketing) and dimensionality reduction (e.g., simplifying complex datasets for visualization) are central. Applications include anomaly detection in cybersecurity, recommendation systems, and topic modeling in large text corpora, where the underlying structure of the data is unknown and must be discovered.

Reinforcement learning is a distinct paradigm where an agent interacts with an environment, learning by receiving feedback in the form of rewards or penalties. The agent's goal is to maximize cumulative rewards over time, learning optimal strategies through trial and error. This approach is especially effective in sequential decision-making scenarios, such as training AI to play games like AlphaGo, autonomous vehicle navigation, robotic control, and algorithmic trading in finance.

Question 2: Describe the architecture and functioning of Convolutional Neural Networks (CNNs) and explain why they are particularly effective for image recognition tasks.

Answer:

Convolutional Neural Networks (CNNs) are specialized deep learning models designed to process grid-like data, particularly images. Their architecture consists of convolutional layers that apply learnable filters to scan the input, extracting features such as edges, textures, and patterns. Pooling layers follow, reducing the spatial dimensions of the data and making the model more efficient while retaining essential information. Activation functions like ReLU introduce non-linearity, and fully connected layers at the end perform classification or regression.

CNNs are exceptionally effective for image recognition because their local connectivity and parameter sharing mirror the spatial structure of images. Early layers detect simple features, while deeper layers combine these to recognize complex objects. Pooling operations provide translation invariance, allowing the network to identify objects regardless of their position in the image. This hierarchical feature

extraction, combined with efficient parameter usage, enables CNNs to outperform traditional neural networks in tasks such as image classification, object detection, semantic segmentation, and facial recognition.

Question 3: Discuss the ethical considerations and potential societal impacts of implementing artificial intelligence systems in critical decision-making processes.

Answer:

The use of AI in critical decision-making brings significant ethical and societal challenges. Algorithmic bias is a primary concern, as AI systems trained on historical data can perpetuate or even amplify existing prejudices, leading to unfair outcomes in areas like hiring, lending, healthcare, and criminal justice. For example, facial recognition systems have demonstrated higher error rates for women and people of color, raising concerns about discrimination and equity.

Transparency and explainability are also crucial. Many advanced AI models, especially deep neural networks, operate as "black boxes," making it difficult to understand or explain their decisions. This lack of interpretability can undermine trust, particularly in high-stakes domains such as healthcare, where understanding the rationale behind recommendations is essential for safety and accountability.

Privacy is another significant concern, as AI systems often require access to large amounts of personal data, raising questions about consent, data security, and surveillance. Accountability is also complex, as responsibility for AI-driven decisions may be distributed among data collectors, developers, system integrators, and end-users, making it difficult to assign liability when harm occurs.

Additionally, the widespread adoption of AI can disrupt labor markets, potentially leading to job displacement and increased economic inequality if the benefits of automation are not broadly shared. Addressing these challenges requires a combination of technical solutions (such as fairness-aware algorithms and explainable AI), regulatory oversight, and ongoing stakeholder engagement to ensure AI systems are ethical, transparent, and aligned with societal values.

Question 4: Explain the concept of transfer learning in deep neural networks and discuss its advantages and limitations.

Answer:

Transfer learning is a technique where a model developed for one task is adapted for a different, but related, task. In deep learning, this typically involves using a neural network pre-trained on a large dataset (such as ImageNet for images) and fine-tuning it for a new task with less data. The process usually includes replacing the final layers of the pre-trained model with new layers suited to the target task and retraining either just these new layers or the entire network with a lower learning rate.

The main advantage of transfer learning is that it leverages knowledge from large, diverse datasets, making it possible to train effective models even when the target dataset is small. This reduces the need for extensive labeled data, speeds up training,

and often leads to better performance by providing a strong starting point for learning. Transfer learning also helps prevent overfitting by transferring general feature representations learned from the source task.

However, transfer learning is most effective when the source and target tasks are similar; if the domains are too different, the transferred knowledge may not be useful or could even hinder performance (a phenomenon known as negative transfer). Additionally, pre-trained models may carry over biases from their original training data, and their architectures may not always be optimal for the new task.

Question 5: Describe the principles of natural language processing (NLP) and how transformer-based models like BERT have revolutionized language understanding tasks.

Answer:

Natural Language Processing (NLP) is a field focused on enabling computers to analyze, interpret, and generate human language. Traditional NLP relied on rule-based systems and statistical models, which often struggled to capture the complexity and context of language.

Transformer-based models, such as BERT (Bidirectional Encoder Representations from Transformers), have revolutionized NLP. Transformers use self-attention mechanisms, allowing the model to consider relationships between all words in a sentence simultaneously, rather than processing them sequentially. BERT introduced bidirectional context, enabling the model to understand the meaning of a word based on both its left and right surroundings.

These models are pre-trained on massive text corpora using tasks like masked language modeling and next sentence prediction, then fine-tuned for specific applications such as sentiment analysis, question answering, and text classification. This approach allows for the development of rich, contextual word representations that capture both semantic and syntactic nuances. As a result, transformer-based models have set new standards for NLP performance, enabling more accurate and natural human-computer interactions. Despite challenges such as high computational requirements and potential biases, transformers have become the backbone of modern NLP systems.